

AI Workflow Skill Map

Choose a practice by the problem you have—not by a slash-command menu. Commands are optional examples; judgment, evidence, privacy, and human approval are the learning objectives.

STUDENT REFERENCE

VENDOR-NEUTRAL FIRST

STATIC-PROJECT BOUNDARY

UPDATED AUGUST 20, 2026

Start with the lowest-power move that can answer the question. A conversation may be enough. Use a file-writing, publishing, committing, account-connected, or secret-handling workflow only when the task actually requires it and the instructor has approved the boundary.

1. Route by situation

Clarify
Ask for plain language, the missing premise, and one concrete example. Restate the idea yourself.
STOP: CLARITY IS NOT VERIFICATION

Decide
Surface assumptions, alternatives, criteria, and tradeoffs. Humans retain decisions and approval.
STOP: DO NOT LET AI INVENT THE OWNER'S ANSWER

Investigate
Find where the answer lives: source material, another person, or a small design probe.
STOP: LABEL UNKNOWNNS AND WEAK EVIDENCE

Build
Write a small behavior specification, slice vertically, and build the smallest reviewable path.
STOP: UNCLEAR BEHAVIOR RETURNS TO DECIDE

Verify
Define expected behavior, reproduce failures, test observable outcomes, review changes, and triage findings.
STOP: REDACT SENSITIVE EVIDENCE

Transfer
Create a handoff only when work must move to another person, harness, directory, repository, or side task.
STOP: DO NOT COPY SECRETS OR PRIVATE DATA

CLARIFY → LOCATE ANSWER → DEFINE → BUILD SMALL → VERIFY →

APPROVE OR REVISE

2. Where does the missing answer live?

THE ANSWER IS IN...	USE THIS PRACTICE	STUDENT ARTIFACT	DO NOT...
Your own judgment	Requirements interview or structured self-questioning	Decision, criteria, and one rejected alternative	Agree automatically with the model's recommendation
Course or project material	Source-bound reading and document critique	Quoted location, careful inference, and open question	Treat a summary as stronger than the original source

THE ANSWER IS IN...	USE THIS PRACTICE	STUDENT ARTIFACT	DO NOT...
One other person's knowledge	Single-recipient discovery questionnaire	Purpose, context, priority questions, and answer stubs	Ask AI to guess policy, preference, or private facts
Nobody's knowledge yet	Disposable design probe	Assumption, observation, and unresolved question	Present a mockup as production or research evidence
Observable system behavior	Reproduction, test case, and controlled debugging	Expected versus actual result and verification record	Change several variables or expose private logs

3. Match authority to risk

Read, explain, or draft locally

Usually low-risk with allowed material. Still verify factual claims and disclose submitted assistance.

Write or change project files

Review the exact target, preserve unrelated work, inspect the diff, and test before acceptance.

Publish, commit, connect, or handle secrets

Requires explicit scope and approval. Use instructor guidance; never place credentials or private data in prompts.

No skill installed? Describe the practice in ordinary language. Example: "Ask me one question at a time until the user, boundary, success test, and main risk are clear. Recommend an answer, but wait for my decision." Grading should measure the quality of the reasoning and artifact, not whether a command ran.

APPENDIX

AIHero 25-Skill Course Audit

This appendix records how the August 20, 2026 AIHero collection maps to LearnAI. It is an adaptation review, not an endorsement or independent validation. Skill behavior may change; open the current source page before installing or relying on a command.

CORE PRACTICE

OPTIONAL

INSTRUCTOR

NOT FOR COURSE

#	AIHERO SKILL	LEARNAI TREATMENT	WHY OR PLACEMENT	STATE / RISK NOTE
1	/setup-matt-pocock-skills	INSTRUCTOR	Repository setup is unnecessary for required student work.	Writes agent documentation and tracker configuration.
2	/ask-matt	REFERENCE ONLY	Use the vendor-neutral situation map instead of teaching a vendor router.	May route using incomplete or stale skill summaries.
3	/grill-with-docs	INSTRUCTOR	Advanced Week 8/10 document-grounded planning.	May write CONTEXT and architecture-decision files.
4	/to-spec	CORE PRACTICE	Week 10: turn settled decisions into a small behavior specification.	The original workflow may publish an issue; LearnAI uses a local text artifact.
5	/to-tickets	CORE PRACTICE	Week 10: create small vertical slices with blocker notes.	The original workflow publishes tracker tickets; LearnAI uses local planning.
6	/implement	INSTRUCTOR DEMO	Show only on a tiny controlled example.	Can edit, test, review, and commit on the current branch.
7	/code-review	CORE CONCEPT	Weeks 10/13: separate "built correctly?" from "built the right thing?"	Command expects a fixed diff and specification; may use subagents.
8	/wayfinder	NOT FOR COURSE	Too much multi-session issue-map overhead for a beginner static project.	Creates and maintains tracker artifacts that can become stale.
9	/prototype	CORE PRACTICE	Week 8: a disposable design probe for one unresolved question.	Original flow may create a branch; LearnAI uses paper or single-file probes.
10	/research	EXISTING PRACTICE	Week 11 already requires primary sources, lateral reading, and claim-level support.	Original skill writes a research artifact and may use costly background agents.
11	/improve-codebase-architecture	INSTRUCTOR	Advanced architecture enrichment only.	Creates a report and may trigger stateful design discussions.
12	/diagnosing-bugs	CORE PRACTICE	Week 9: reproduce, minimize, change one variable, and verify.	Logs or diagnostic artifacts can expose credentials or private data.
13	/resolving-merge-conflicts	SAFETY CONCEPT	Teach what a conflict is and when to stop and ask for help.	The command resolves and commits; beginners should not automate this.

#	AIHERO SKILL	LEARNAI TREATMENT	WHY OR PLACEMENT	STATE / RISK NOTE
14	/triage	EXISTING PRACTICE	Week 13 already classifies findings before fixing or documenting them.	Original workflow may label, comment, or write tracker artifacts.
15	/wizard	NOT FOR COURSE	Outside the static, no-secrets course boundary.	Generates shell guidance for dashboards, environment variables, and secrets.
16	/grill-me	EXISTING PRACTICE	Weeks 8/10: stress-test requirements while humans retain decisions.	Chat-only, but a model can still over-recommend or over-agree.
17	/handoff	EXISTING PRACTICE	Week 10: portable brief only when work must travel.	Writes a temporary file; temporary storage is not durable.
18	/to-questionnaire	EXISTING PRACTICE	Week 8: one recipient who holds missing knowledge.	Writes local Markdown but does not send it; review before sharing.
19	/teach	OPTIONAL	Week 11 self-study experiment, not a substitute for course teaching.	Creates a multi-file teaching workspace; generated content and quizzes need review.
20	/wait-what	CORE PRACTICE	Week 2: repair an explanation with plain language, a premise, and an example.	Chat-only; clearer output is not necessarily more accurate.
21	/writing-for-agents	OPTIONAL	Week 12: concise, staged, completion-aware meta-instructions.	Instructions do not guarantee compliance; verify actions and outputs.
22	/codebase-design	INSTRUCTOR	Specialized software-design vocabulary beyond the beginner objective.	A reference without a driver can trigger unintended exploration.
23	/domain-modeling	OPTIONAL	Capstone glossary exercise when terms are genuinely ambiguous.	Writes CONTEXT and may offer architecture-decision records.
24	/grilling	NO SEPARATE LESSON	The interview mechanism is already adapted through requirements activities.	Model-invoked reference may use subagents for factual questions.
25	/tdd	CORE PRACTICE LITE	Week 9: define one observable failing case before a fix when useful.	The full engineering method needs agreed test boundaries and stronger code knowledge.

Installation is not a course requirement. Before installing any third-party skill, inspect its current source, requested permissions, file-writing behavior, external connections, update mechanism, and license. A useful description does not prove that every model or environment will follow it reliably.

Primary collection reviewed: **AIHero Skills**, 25 linked pages as presented on August 20, 2026. This LearnAI map paraphrases workflow ideas for a beginner educational setting and adds local privacy, evidence, assessment, static-project, and human-approval boundaries.